



Large Scale Machine Learning

TOTAL POINTS 5

1. Suppose you are training a logistic regression classifier using stochastic gradient descent. You find that the cost (say, $\text{cost}(\theta, (x^{(i)}, y^{(i)}))$), averaged over the last 500 examples, plotted as a function of the number of iterations, is slowly increasing over time. Which of the following changes are likely to help?

1 point

- ☐ Try using a larger learning rate α .
- ☐ Try averaging the cost over a larger number of examples (say 1000 examples instead of 500) in the plot.
- ☒ Try using a smaller learning rate α .
- ☐ This is not an issue, as we expect this to occur with stochastic gradient descent.



2. Which of the following statements about stochastic gradient

1 point

descent are true? Check all that apply.

- ☒ If you have a huge training set, then stochastic gradient descent may be much faster than batch gradient descent.
- ☐ One of the advantages of stochastic gradient descent is that it uses parallelization and thus runs much faster than batch gradient descent.
- ☒ Before running stochastic gradient descent, you should randomly shuffle (reorder) the training set.
- ☐ In order to make sure stochastic gradient descent is converging, we typically compute $J_{\text{train}}(\theta)$ after each iteration (and plot it) in order to make sure that the cost function is generally decreasing.



3. Which of the following statements about online learning are true? Check all that apply.

1 point

- ☒ One of the advantages of online learning is that if the function we're modeling changes over time (such as if we are modeling the probability of users clicking on different URLs, and user tastes/preferences are changing over time), the online learning algorithm will automatically adapt to these changes.
- ☐ Online learning algorithms are most appropriate when we have a fixed training set of size m that we want to train on.
- ☐ When using online learning, you must save every new training example you get, as you will need to reuse past examples to re-train the model even after you get new training examples in the future.
- ☒ Online learning algorithms are usually best suited to problems where we have a continuous/non-stop stream of data that we want to learn from.



4. Assuming that you have a very large training set, which of the

1 point

following algorithms do you think can be parallelized using

map-reduce and splitting the training set across different

machines? Check all that apply.

- ☐ An online learning setting, where you repeatedly get a single example (x, y) , and want to learn from that single example before moving on.
- ☐ Logistic regression trained using stochastic gradient descent.
- ☒ A neural network trained using batch gradient descent.



☒ Linear regression trained using batch gradient descent.

5. Which of the following statements about map-reduce are true? Check all that apply.

1 point

- ☒ If you have only 1 computer with 1 computing core, then map-reduce is unlikely to help.
- ☐ If we run map-reduce using N computers, then we will always get at least an N -fold speedup compared to using 1 computer.
- ☒ When using map-reduce with gradient descent, we usually use a single machine that accumulates the gradients from each of the map-reduce machines, in order to compute the parameter update for that iteration.
- ☒ Because of network latency and other overhead associated with map-reduce, if we run map-reduce using N computers, we might get less than an N -fold speedup compared to using 1 computer.

☐ I understand that submitting work that isn't my own may result in permanent failure of this course or deactivation of my Coursera account.



[Learn more about Coursera's Honor Code](#)

Lucas

Save

Submit

